

Problem 4

a)

```
# truncated power base function  
# when added to model for a cubic polynomial this  
# keeps f continuous for 1st and 2nd derivatives at each knot  
# x is predictor  
# z is knot  
  
h <- function(x,z) {  
  ifelse(x>z,(x-z)^3,0)  
}
```

b)

```
#  
# xs is vector of predictors  
# h is a matrix of truncated power base functions  
  
hs <- function(xs,z) {  
  sapply(xs,h, z=z)  
}
```

c)

```
#  
# xs is predictors  
# zs is knots  
  
splinebasis <- function(xs,zs) {  
  mhs <- matrix(data=0,nrow=length(xs),ncol=length(zs))  
  for(j in 1:length(zs)){  
    mhs[,j] = sapply( xs, hs, zs[j])  
  }  
  m <- cbind(xs, xs^2, xs^3, mhs)  
}
```

d)

```
set.seed(2019)  
x = runif(100)  
y = sin(10*x)
```

e)

```
knots=c(1/4,2/4,3/4)  
  
out3 = lm(y~I(splinebasis(x,knots)))  
  
print(out3)
```

```
##
```

```
## Call:
## lm(formula = y ~ I(splinebasis(x, knots)))
##
## Coefficients:
##              (Intercept)  I(splinebasis(x, knots))xs
##                -0.1197                16.2295
##   I(splinebasis(x, knots))  I(splinebasis(x, knots))
##                -68.1331                56.4607
##   I(splinebasis(x, knots))  I(splinebasis(x, knots))
##                62.9915                -284.9395
##   I(splinebasis(x, knots))
##                259.4466
```


```
yhat <- predict(out3)

reorder <- order(x)

plot(x,y)

lines(x[reorder],y[reorder])

points(x,yhat, col='green')
lines(x[reorder],yhat[reorder], col='green')
```



The spline (in green) does an OK job of approximating the function. The green line is fairly close to the black line but the spline is a bit low both on the left maximum and the central minimum.

f)

```
k1=c(1/2)
k2=c(1/3,2/3)
k3=c(1/4,2/4,3/4)
k4=c(1/5,2/5,3/5,4/5)
k5=c(1/6,2/6,3/6,4/6,5/6)
```


```

k6=c(1/7,2/7,3/7,4/7,5/7,6/7)
k7=c(1/8,2/8,3/8,4/8,5/8,6/8,7/8)
k8=c(1/9,2/9,3/9,4/9,5/9,6/9,7/9,8/9)
k9=c(1/10,2/10,3/10,4/10,5/10,6/10,7/10,8/10,9/10)


runner <- function (k) {
  m <- lm(y~I(splinebasis(x,k)))
  yhat <- predict(m)
  reorder <- order(x)
  plot(x,y)
  lines(x[reorder],y[reorder])
  points(x,yhat, col='green',pch=5)
  lines(x[reorder],yhat[reorder], col='green',lty=2)
}

```


```
o1 <- runner(k1)
```




```
o2 <- runner(k2)
```




```
o3 <- runner(k3)
```




```
o4 <- runner(k4)
```




```
o5 <- runner(k5)
```




```
o6 <- runner(k6)
```




```
o7 <- runner(k7)
```



```
o8 <- runner(k8)
```



```
o9 <- runner(k9)
```



g)

Yes, as the number of knots increases we expect the curve to become a better approximation of the true curve because it is increasingly flexible and better able to match the true data. Improvement will slow and then stop as the number DoF in the model approach the number of unique x values. However, in the real world when the true data is unknown, increasing the knots would likely lead to overfitting. Picking a good k could be done by using a validation or cross-validation approach.